

CCAO-F

Pre-Exam Recap

Claude Certified Associate - Foundations

Seven domains. One reference. Read it the night before.

1

2

3

4

5

6

7

The exam at a glance

120

minutes

60

questions, MCQ and
multiple response

720+

passing score
(100 to 1,000 scale)

12

months validity

Who it is for. Business users of Claude: project managers, business analysts, delivery leads. No coding required.

What it really tests. Judgment, not memorization. Every question is a workplace scenario with four options, and usually two of them look right. The skill being measured is choosing the defensible one.

Delivery is proctored, online or at a test center. Format details are as of this deck's date - always verify the current exam guide on Anthropic's official certification page before your sitting.

The blueprint - where the marks are



Orange rows carry more than a third of the exam between them. If time is short tonight, reread Domains 2, 4, and 6.

1

Prompting and Task Execution

14%

THE MOVES

1

Say it clearly. Clear, explicit, specific. Name the role, context, audience, format, and what to skip.

2

Structure it. Fence instructions, context, and input with tags. Number ordered steps. Long document at the top, question at the end.

3

Split and iterate. Chain big jobs so one output feeds the next. Draft, review, refine. Test against written criteria, not feel.

4

Match the task. Research wants cross-checking. Deep analysis wants thinking room. Formatting wants three to five varied examples.

REMEMBER

- The colleague test: if a colleague with no context is confused by your prompt, Claude is too.
- Phrase rules as do, not do not - and show the wanted shape instead of banning the unwanted one.
- Not every failure is a prompt problem. Good but slow or costly means a different model, not rewrite six.
- Formatting inconsistency is fixed with examples, not with more instructions.

THE TRAP

The answer that adds a negative instruction ("do not change the structure") without ever showing the wanted structure. It feels like an instruction. It is a hope.

2

Output Evaluation - Part A: can you trust it?

21%

THE MOVES

1

Judge it. Score against your checklist - faithful to source, complete, consistent, right tone - written before you read the output.

2

Spot the four tells. A citation that leads nowhere. A precise number with no source. An answer that contradicts itself. Confidence with nothing behind it.

3

Make it provable. Allow "I don't know" (the single most effective line). Ask for quotes first. Require a citation per claim - withdraw what can't be cited.

4

Audit on contact. One real error found means others are likely. Finding it starts the audit, not ends it.

REMEMBER

- Confidence and accuracy are unrelated. Fluent is not faithful.
- A citation is not verification. Verification means you open the source and read it.
- Self-reported certainty is never a signal. The checker cannot be the thing being checked.
- For big batches, sample and estimate the error rate - never skim and trust.

THE TRAP

"Ask Claude to double-check its own answer." It feels diligent. It is the same model marking its own homework - the exam's favorite wrong answer.

2

Output Evaluation - Part B: who signs it, how it ships

21%

THE MOVES

- 1 Gate high stakes.** Rights, health, money, or livelihood - a qualified human reviews before anything acts. Legal, medical, insurance, finance, hiring, testing, media.
- 2 Disclose.** People are told they are dealing with AI output. Consumer-facing AI always discloses.
- 3 Fit the reader.** Name the reader in one phrase, then shape it: executives get the point, practitioners get the steps, partners get prose.
- 4 Pick the container.** Reusable work in a document. Data for systems in a fixed structure. One-off answers as plain prose.

REMEMBER

- A strict format guarantees the shape, never the facts. Fields present, totals still wrong.
- Prompt style leaks into output style - a bulleted prompt comes back as bullets.
- If nobody signed it, it didn't ship. Keep the approval where you can point at it.
- Rejections and adverse decisions are where the human-review rule applies hardest.

THE TRAP

Any option where the model certifies its own high-stakes output, or the sign-off is skipped because "the reasoning looks clean." Presentation is not review.

3

Product and Model Selection

12%

THE MOVES

- 1 Pick the door by the verb.** Gather = Research. Compute = code execution. Recurring = Projects. Reusable deliverable = Artifacts. Repeated procedure = Skills. Remember = Memory.
- 2 Know the ladder.** Haiku for volume, Sonnet for balance, Opus and Fable for the deepest work. Capability up means cost up and speed down - nobody gets all three.
- 3 Tune before you switch.** Missing info? Supply it. Still shallow? More thinking on the same model. Still short? Climb one tier.
- 4 Mind the window.** Context is finite and quality sags as it fills; the middle gets missed. Prune, then compact, then restart with key state.

REMEMBER

- Fit the rung to the task, measured on your own work - never chosen by habit, mood, or the open tab.
- Retry without new information changes nothing. Same input, same model, same answer.
- Turning features off does not reduce cost meaningfully - fitting the right tier does.
- Keep decisions, not transcripts. Smaller, focused context is better and cheaper.

THE TRAP

"Run it again." The reflex everyone has. Without new information or new settings, a retry is a lottery ticket, not a lever.

4

Workflow Integration and Solution Design

16%

THE MOVES

1

Find where it pays. Walk the process, mark each step for time eaten and language work. High on both = your step. Rank by impact against effort.

2

Simplest pattern first. Direct prompt, then chain, route, or parallelize. Fixed workflows are predictable; free-running agents are flexible and harder to control - reach for them last.

3

Augment before redesign. If the workflow works, wire Claude into the existing tools via connectors. Redesign only what is already broken.

4

Sell it honestly. Quantify from your own pilot. State cost, failure modes, and which steps keep a human signature. Frame as capacity moved, not people replaced.

REMEMBER

- The most successful implementations start simple and add complexity only when the need is proven.
- Combining Claude with existing tools usually beats rebuilding the workflow from scratch.
- Route easy items to the fast tier, hard ones up the ladder.
- Never borrow a vendor's ROI numbers - measure your own.

THE TRAP

"Rebuild it end to end around an autonomous agent." It sounds like vision. It is complexity before evidence - and the exam knows it tempts you.

5

Configuration and Knowledge Management

12%

THE MOVES

- 1 Projects = instructions + knowledge.** The quality of both determines the quality of every answer inside. Split projects by purpose - never one overloaded workspace.
- 2 Restrict the source.** Uploading files is not enough. Say: answer only from these materials, and say unknown for anything not in them.
- 3 Structure the instructions.** Role, constraints, and output format as separate tagged blocks. One long paragraph gets its middle ignored. State priority when rules can clash.
- 4 Maintain it.** Audit instructions, refresh stale files and sync sources, keep examples current, store settled decisions to memory - not chat history.

REMEMBER

- Narrow every connector to the minimum scope. Connecting a whole drive and asking Claude to ignore the noise does not work.
- Static uploads freeze at import; sync sources stay current - know which one each source is.
- Contradictory materials in one project destabilize every answer it gives.
- Memory holds facts and preferences across sessions, never transcripts.

THE TRAP

One giant project holding everything, a whole shared drive connected "to be safe," and instructions in a single paragraph. Each feels convenient. Each degrades accuracy.

6

Governance, Risk, and Responsible Use

15%

THE MOVES

- 1 Know the seven high-risk areas.** Legal, healthcare, insurance, finance, employment and housing, academic testing, media. Human-in-the-loop plus disclosure, always.
- 2 Check fairness.** Is one side over-represented? Are opposing views dismissed? Is a group treated unfairly? Build the habit on every sensitive output.
- 3 Protect the data.** PII anonymized or explicitly authorized. Org policy decides what may enter any AI tool. Ambiguous cases escalate to compliance, not to a guess.
- 4 Communicate limits.** Value with real numbers, limits said out loud: cost, failure modes, the standing need for human oversight.

REMEMBER

- In high-risk work, Claude must never self-verify - a qualified professional signs.
- Hiding AI involvement from consumers violates the usage policy. Disclose at the start.
- Regulated domains have record-keeping duties - confirm outputs can be logged and retained.
- An answer that ignores company policy is wrong even when the technology in it is right.

THE TRAP

Any option that skips disclosure, skips the qualified reviewer, or treats "Claude is confident" as governance. Compliance beats clever, every time.

7

Troubleshooting and Optimization

10%

THE MOVES

- 1 Diagnose in order.** 1 Prompt unclear? Rewrite it. 2 Information missing? Supply it. 3 Reasoning shallow? More effort, then a higher tier. 4 Context stale? Prune, compact, or restart.
- 2 Pull levers by impact.** Prompt first, effort second, model third, context fourth, tools last.
- 3 Validate with criteria.** Define success up front and measure. Many rough automated checks beat a few perfect manual ones.
- 4 Optimize cost by fit.** Right tier for the workload, focused context, features matched to the job - measure, adjust, remeasure.

REMEMBER

- Root cause first, fix second. Retrying without a diagnosis is the domain's named anti-pattern.
- "Feels wrong" is not a finding - compare against written criteria.
- Always-the-top-model wastes budget; features-off-to-save-money fixes nothing.
- Longer context does not mean better answers. Usually the opposite.

THE TRAP

Fixing the symptom you noticed instead of the cause you diagnosed - the exam describes a failure and offers four fixes, only one of which matches the actual root cause.

Answer anatomy - the five traps, exam-wide

THE SELF-CHECK

The model confirms or reviews its own answer. Feels diligent; the checker is the thing being checked.

THE RETRY

Run it again with nothing new. Feels persistent; changes nothing.

THE REBUILD

Replace the whole working process. Feels visionary; it is complexity before evidence.

THE NEGATIVE

"Do not do X" with no example of the wanted shape. Feels like an instruction; it is a hope.

THE CLEVER

Technically smart, policy-blind. Feels efficient; compliance outranks clever.

The trap is the option that feels the most responsible or the most impressive. When an answer flatters you, slow down.

Every practice question in the series ends on one of these. So does the real thing.

Eight rules that answer half the paper

1

Confidence is not accuracy.

2

A citation is not verification - open the source.

3

High stakes means a human signs.

4

Claude drafts and synthesizes. You decide.

5

Simplest thing first. Complexity must earn its place.

6

Missing info, then effort, then escalate. Never a blind retry.

7

Org policy outranks the clever answer.

8

Restrict the source. Narrow the scope.

If two options survive elimination, the one that obeys more of these rules wins.

In the room - technique

1 Eliminate first

Two of four options usually die on sight - the do-nothing and the do-everything. Kill them before you think hard.

2 Read for stakes words

Hiring, lending, diagnosis, applicants, patients - the moment a person's rights, health, money, or job is touched, the human-review answer is favored.

3 Read for verbs

Gather, compute, reuse, repeat, remember - task verbs point at the feature the question wants named.

4 Then apply a rule

Between the two survivors, the golden rules decide. The trap is the one that feels responsible or impressive.

Pace: sixty questions in one hundred twenty minutes is two minutes each. Elimination buys you the time; flag and return rather than stall.

The night before



Scan the blueprint - know where the marks are (Domains 2, 4, 6 are the heavy rows).



Reread the five traps and the eight rules. They are the exam's skeleton.



Verify the current format on the official certification page - details can change.



Rewrite one real work prompt with a role, tagged sections, and one example. Ten minutes of practice beats an hour of reading.



Sleep. Judgment is the thing being tested, and judgment needs rest.

Good luck. You saw every trap coming.